Genetic Variants Conflict Prediction Using Machine Learning Algorithms

The advent of next-generation sequencing (NGS) technologies has revolutionized the field of genomics, enabling researchers to delve deeper into the genetic underpinnings of human health and disease. These high-throughput sequencing methods generate vast amounts of genetic data, including information about genetic variations, such as single nucleotide polymorphisms (SNPs) and insertions/deletions (indels). While these variations contribute to the unique genetic makeup of individuals and populations, they also pose significant challenges for interpreting their clinical significance and impact on disease susceptibility and progression.

One of the most critical challenges in genetic variant interpretation arises from the presence of conflicting interpretations across different databases and prediction algorithms. These conflicts can stem from various factors, including differences in underlying data sources, methodologies, and variant classification criteria. Such discrepancies create uncertainty for clinicians and researchers, hindering the effective translation of genetic information into clinical practice and personalized medicine approaches.

The Challenge of Conflicting Interpretations:

Several factors contribute to the challenge of conflicting genetic variant interpretations:

Heterogeneity of Data Sources: Different databases and prediction algorithms may rely on varying data sources, such as population-specific genetic studies, functional assays, or in silico predictions. This heterogeneity can lead to inconsistencies in variant classifications, as the evidence supporting pathogenicity may differ across sources.

Variability in Methodologies: Various prediction algorithms employ diverse methodologies, including sequence conservation analysis, protein structure prediction, and machine learning approaches. Each method has its strengths and limitations, leading to potential discrepancies in variant classifications.

Evolving Classification Criteria: The criteria for classifying genetic variants as pathogenic, likely pathogenic, benign, likely benign, or variants of uncertain significance (VUS) are continuously evolving as new knowledge emerges. This can create challenges in harmonizing classifications across different databases and algorithms.

Limited Functional Data: For many genetic variants, particularly rare ones, functional data from experimental studies may be scarce. This lack of evidence can lead to uncertainty and conflicting interpretations regarding the variant's impact on protein function and disease risk.

The Potential of Machine Learning:

Machine learning (ML) presents a promising approach to address the challenge of conflicting genetic variant interpretations. ML algorithms excel at identifying complex patterns and relationships within large datasets. By leveraging this capability, researchers can develop models that integrate diverse data sources, including genetic, clinical, and functional information, to improve the accuracy and consistency of variant classifications.

Several ML algorithms have been explored for genetic variant classification, including:

Support Vector Machines (SVMs): SVMs are powerful algorithms that can effectively classify data points based on their features. In the context of variant classification, SVMs can be trained on known pathogenic and benign variants to learn the distinguishing features between the two classes and subsequently classify novel variants.

Random Forests (RFs): RFs are ensemble learning methods that combine multiple decision trees to improve prediction accuracy and reduce overfitting. They are particularly well-suited for handling high-dimensional data with complex interactions among features.

Deep Learning (DL): DL models, such as convolutional neural networks (CNNs) and recurrent neural networks (RNNs), have shown exceptional performance in various classification tasks. In the context of genetic variant classification, DL models can learn intricate patterns from sequence data, protein structure information, and other relevant features.

Research Aims and Objectives:

This research aims to develop and evaluate novel machine learning approaches for classifying conflicting genetic variants. The specific objectives include:

1. Obtain a comprehensive dataset of genetic variants with conflicting interpretations: This dataset will serve as the foundation for training and evaluating the machine learning models.

2. Exploring and comparing different machine learning algorithms: A range of ML algorithms, including SVMs, RFs, and DL models, will be investigated to identify the most effective approach for classifying conflicting variants.

3. Developing interpretable models: The research will focus on developing models that provide insights into the features and decision-making process behind their classifications. This will enhance trust and understanding of the models' predictions.

4. Evaluating model performance and generalizability: The performance of the developed models will be rigorously evaluated on a held out dataset to ensure their generalizability and applicability to diverse populations and variant types.

Potential Impact and Applications:

The successful development of accurate and reliable methods for classifying conflicting genetic variants has the potential to significantly impact various areas of genomics and healthcare:

Improved clinical decision-making: By providing more accurate and consistent variant classifications, this research can empower clinicians to make informed decisions regarding patient care, including risk assessment, diagnosis, and treatment planning.

Enhanced genetic counseling: Accurate variant classification can aid genetic counselors in providing comprehensive and reliable information to patients and families about the potential implications of genetic variants.

Advancements in personalized medicine: The research can contribute to the development of personalized medicine approaches by enabling the identification of individuals who may benefit from targeted therapies or preventive measures based on their genetic profiles.

Facilitating research discoveries: Accurate variant classification is crucial for advancing research into the genetic basis of diseases and identifying potential therapeutic targets.

This research holds promise for addressing the challenge of conflicting genetic variant interpretations, leading to more accurate diagnoses, improved patient care, and advancements in personalized medicine and genomic research. By harnessing the power of machine learning, we can unlock the full potential of genetic information for improving human health and well-being.

Methodology:

1. Data Collection:

The data for this research paper was downloaded from Kaggle, a popular platform for data science competitions, datasets, and projects. Kaggle provides a diverse range of datasets across various domains, allowing researchers and data scientists to access, analyze, and derive insights from real-world data.

2. Understanding the Data Structure:

Prior to data analysis, I reviewed the data documentation and explored the dataset's structure, including the variables, data types, and number of missing, duplicated and unique values.

3. Exploratory Data Analysis (EDA):

EDA techniques were applied to gain a deeper understanding of the dataset, identify patterns, trends, and outliers, and lay the foundation for subsequent analysis in the research paper.

4. Data Preprocessing:

The collected dataset underwent preprocessing steps to clean, transform, and prepare the data for analysis. This included handling missing values, splitting the dataset, encoding categorical variables, and scaling.

5. Model Selection and Training:

- Identify suitable machine learning algorithms for conflicting genetics variants analysis.

- Experiment with classification models such as Random Forest, Support Vector Machines, Gradient Boosting, and Neural Networks.

- Split the dataset into training and validation sets for model training and evaluation.

6. Hyperparameter Tuning:

- Utilize techniques like grid search or random search to optimize model hyperparameters.

- Perform cross-validation to ensure the robustness and generalization of the model.

7. Model Evaluation:

- Assess the performance of the trained models using evaluation metrics like accuracy, precision, recall, F1-score, and area under the ROC curve.

- Analyze model predictions and investigate the impact of machine learning on resolving conflicting genetic variants.

8. Validation and Interpretation:

- Validate the machine learning model on independent datasets to assess its generalization capability.

- Interpret the results, analyze the model's decision-making process, and identify key factors influencing conflicting variant resolution.

9. Implementation and Deployment:

- Develop a scalable and user-friendly tool based on the trained machine learning model to automate the resolution of conflicting genetic variants.

- Ensure the tool's usability and accessibility for geneticists, clinicians, and researchers in the field of precision medicine.

By following this methodology, we aim to leverage machine learning techniques to effectively address conflicting genetics variants, improve variant classification accuracy, and contribute to advancing the field of precision genetics analysis.

Exploring Machine Learning Algorithms for Genetic Variant Classification

Machine learning (ML) offers a powerful toolkit for addressing the challenge of conflicting genetic variant interpretations. Various ML algorithms have demonstrated their effectiveness in classifying complex data patterns, making them suitable candidates for this task. This research explores a selection of diverse ML algorithms, each with its unique strengths and limitations:

1. Random Forest (RF):

Approach: Random Forests are ensemble learning methods that combine multiple decision trees to improve prediction accuracy and robustness. Each tree is trained on a random subset of features and data samples, promoting diversity and reducing overfitting. The final prediction is made by aggregating the predictions of individual trees, typically through majority voting for classification tasks.

Strengths: RFs are known for their ability to handle high-dimensional data with complex interactions among features. They are also relatively robust to outliers and noise and require minimal parameter tuning.

Limitations: RFs can be computationally expensive to train, especially with large datasets and a high number of trees. Additionally, their decision-making process can be difficult to interpret, limiting their explainability.

2. Logistic Regression (LR):

Approach: Logistic regression is a statistical method used for binary classification tasks. It models the relationship between input features and the probability of belonging to a specific class using a logistic function. The model learns the weights for each feature that maximize the likelihood of the observed data.

Strengths: LR is a simple and interpretable model, allowing for easy understanding of the contribution of each feature to the prediction. It is also computationally efficient and well-suited for problems with a clear separation between classes.

Limitations: LR assumes a linear relationship between features and the log-odds of the outcome, which may not always hold true for complex datasets. Additionally, it can be sensitive to outliers and multicollinearity among features.

3. Gaussian Naive Bayes (GNB):

Approach: GNB is a probabilistic classification algorithm based on Bayes' theorem with the assumption of independence among features. It calculates the probability of a data point belonging to each class, given its feature values, and assigns the class with the highest probability.

Strengths: GNB is a simple and efficient algorithm that requires minimal training data. It is also relatively insensitive to irrelevant features and noise.

Limitations: The assumption of independence among features rarely holds true in real-world datasets, which can limit the accuracy of GNB.

4. Multilayer Perceptron (MLP):

Approach: MLP is a type of artificial neural network with multiple layers of interconnected nodes or neurons. Each neuron applies a non-linear activation function to its weighted inputs, allowing the network to learn complex patterns and relationships within the data.

Strengths: MLPs are powerful models capable of learning non-linear relationships and achieving high accuracy on complex datasets.

Limitations: MLPs can be prone to overfitting, requiring careful regularization and hyperparameter tuning. Additionally, their decision-making process is often difficult to interpret.

5. K-Nearest Neighbors (KNN):

Approach: KNN is a non-parametric algorithm that classifies data points based on the k nearest neighbors in the training data. The distance between data points is typically measured using Euclidean distance or other distance metrics.

Strengths: KNN is simple to implement and requires no training phase. It is also effective for datasets with non-linear relationships among features.

Limitations: KNN can be computationally expensive during prediction, especially with large datasets. Additionally, its performance is sensitive to the choice of k and the distance metric used.

6. XGBoost:

Approach**:** XGBoost is a gradient boosting algorithm that builds an ensemble of decision trees sequentially. Each tree is trained to correct the errors of the previous trees, resulting in a powerful and accurate model.

Strengths: XGBoost is known for its high accuracy and efficiency. It can handle missing data and is relatively robust to outliers.

Limitations: XGBoost requires careful hyperparameter tuning and can be prone to overfitting if not properly configured.

Choosing the Best Algorithm:

The choice of the most suitable ML algorithm for conflicting genetic variant classification depends on several factors, including the characteristics of the dataset, the desired level of interpretability, and computational resources. This research will employ a comparative approach, evaluating the performance of each algorithm on the curated dataset of conflicting variants. Metrics such as accuracy, precision, recall, F1 score, and area under the ROC curve (AUC) will be used to assess model performance. Additionally, the interpretability of the models will be evaluated to ensure that their predictions can be understood and trusted.

By exploring and comparing these diverse ML algorithms, the research aims to identify the most effective approach for classifying conflicting genetic variants, paving the way for improved clinical decision-making and advancements in personalized medicine and genomic research.

Evaluating Machine Learning Models for Genetic Variant Classification: Performance Metrics and Interpretability

Evaluating the performance of machine learning (ML) models is crucial for assessing their effectiveness and ensuring their suitability for real-world applications. In the context of classifying conflicting genetic variants, several metrics are essential for understanding the strengths and limitations of different models:

1. Accuracy:

Definition: Accuracy measures the proportion of correct predictions made by the model out of the total number of predictions. It is a simple and intuitive metric, often used as a preliminary indicator of model performance.

Formula: Accuracy = (True Positives + True Negatives) / (Total Predictions)

Strengths: Easy to understand and interpret. Provides a general overview of the model's performance.

Limitations: Can be misleading for imbalanced datasets, where one class is significantly more prevalent than others. In such cases, a model that simply predicts the majority class may achieve high accuracy but fail to identify the minority class effectively.

2. Precision:

Definition: Precision measures the proportion of true positive predictions out of all positive predictions made by the model. It indicates the model's ability to avoid false positives, i.e., incorrectly classifying a variant as conflicting when it is actually does not has a conflicting classification.

Formula: Precision = True Positives / (True Positives + False Positives)

Strengths: Useful for applications where minimizing false positives is critical, such as in clinical settings where misclassifying a variant could lead to unnecessary interventions or anxiety for patients.

Limitations: May not be the most informative metric for imbalanced datasets, where the number of true positives is relatively small.

3. Recall (Sensitivity):

Definition: Recall measures the proportion of true positive predictions out of all actual positive cases. It reflects the model's ability to identify all pathogenic variants correctly.

Formula: Recall = True Positives / (True Positives + False Negatives)

Strengths: Important for applications where identifying all true positives is crucial, such as in screening programs or research studies aimed at understanding the genetic basis of diseases.

Limitations: May not be the most informative metric when false positives are also a concern.

4. F1 Score:

Definition: The F1 score is the harmonic mean of precision and recall, providing a balanced measure that considers both false positives and false negatives.

Formula: F1 Score = 2 (Precision Recall) / (Precision + Recall)

Strengths: Provides a single metric that balances precision and recall, making it useful for comparing models with different performance profiles.

Limitations: May not be as intuitive to interpret as accuracy or individual precision and recall values.

5. Area Under the Receiver Operating Characteristic Curve (AUROC):

Definition: The ROC curve plots the true positive rate (sensitivity) against the false positive rate (1 - specificity) at various classification thresholds. The AUROC represents the area under this curve and provides a measure of the model's ability to discriminate between classes across different thresholds.

Strengths: Provides a threshold-independent evaluation of the model's performance, making it suitable for comparing models even if they operate at different classification thresholds.

Limitations: Can be less informative for highly imbalanced datasets where the ROC curve may appear overly optimistic due to the large number of true negatives.

6. Area Under the Precision-Recall Curve (AUPRC):

Definition: The precision-recall curve plots precision against recall at different classification thresholds. The AUPRC represents the area under this curve and is particularly useful for evaluating models on imbalanced datasets.

Strengths: Provides a more informative evaluation for imbalanced datasets compared to AUROC, as it focuses on the model's ability to correctly classify the minority class (e.g., conflicting variants).

Limitations: Can be more challenging to interpret than AUROC, especially for those unfamiliar with precision-recall curves.

Model Interpretability:

Beyond quantitative performance metrics, understanding the reasoning behind a model's predictions is crucial for building trust and facilitating its adoption in clinical and research settings. Model interpretability allows for:

Identifying potential biases: Examining the features that contribute most to the model's predictions can reveal potential biases in the training data or the model itself.

Understanding model limitations: Analyzing the types of errors made by the model can highlight its limitations and areas for improvement.

Building trust and acceptance: Providing explanations for the model's predictions can increase trust and acceptance among clinicians, researchers, and patients.

Techniques for Enhancing Interpretability:

Several techniques can be employed to enhance the interpretability of ML models for genetic variant classification:

Feature importance analysis: This technique identifies the features that contribute most to the model's predictions, providing insights into the factors driving the classifications.

Model-agnostic interpretation methods: These methods, such as LIME (Local Interpretable Model-agnostic Explanations) and SHAP (SHapley Additive exPlanations), provide explanations for individual predictions by approximating the model locally with simpler, interpretable models.

Visualization techniques: Visualizing the decision-making process of the model, such as through decision trees or activation maps in neural networks, can provide insights into how the model arrives at its predictions.

Result:

Exploratory Data Analysis:

The initial exploration of the dataset reveals several key characteristics:

Size and Structure: The dataset is moderately large, containing 65,188 rows (representing individual samples or observations) and 46 columns (representing different features or variables). This suggests a good amount of data for analysis, but not so large as to be unwieldy.

Missing Data: We encounter a challenge of missing values. 30 of the 46 columns contain missing data, which will require careful handling.

Target Variable: The target variable has two classes encoded as 1 for samples with conflicting classification, and 0 for sample that do not have conflicting classification. The first class (conflicting) has a significantly smaller representation with 12,506 samples compared to the second class’s 37,742 samples. This imbalance needs to be considered during modeling and evaluation as it can bias some algorithms.

Feature Types:

'CLNVC', 'IMPACT', 'Feature_type', 'BIOTYPE' are categorical features while 'POS', 'AF_ESP', 'AF_EXAC', 'AF_TGP', 'ORIGIN', 'CLASS', 'STRAND', 'LoFtool', 'CADD_PHRED', 'CADD_RAW' are numeric
.
'REF', 'CLNDISDB', 'CLNDN', 'CLNHGVS', 'SYMBOL', 'Feature', 'EXON', 'cDNA_position', 'CDS_position', 'Protein_position', 'Amino_acids', 'Codons’ features are also categorical but have a high number of unique categories. Such features tend to overfit the model and therefor need to be handled.
'Feature_type', 'BIOTYPE' features appear to have only one unique value. These features might not be informative for the analysis and could potentially be removed.

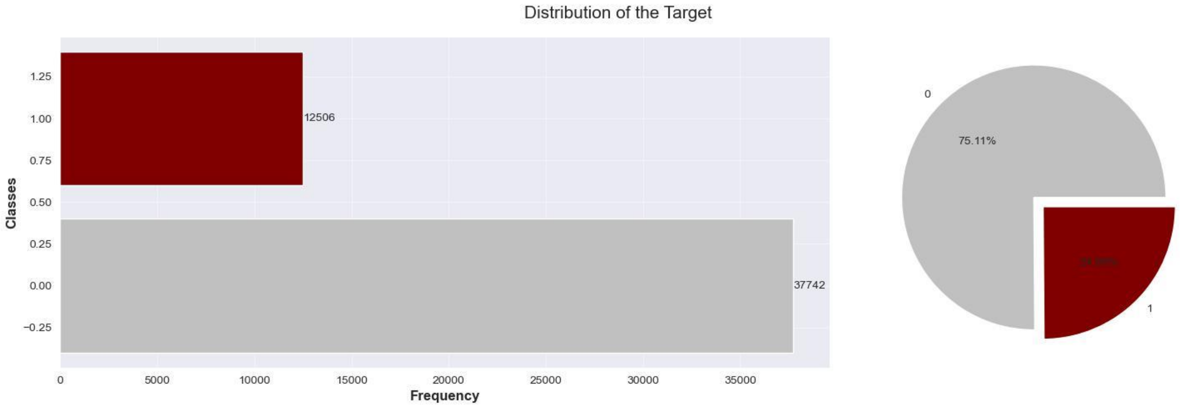

Data Preprocessing

To ensure the quality of our analysis, we embark on a series of preprocessing steps to clean and prepare the data:

Missing Value Management:

Our initial investigation revealed a significant presence of missing values. To address this, we take a two-pronged approach:

Column Removal: We identify columns with more than 50% missing values and deem them unreliable for analysis. These columns are dropped entirely from the dataset.

Row Removal: For remaining columns, we eliminate rows containing missing values. This is a stricter approach, but it ensures we work with complete data and reduces the risk of bias or errors introduced by imputation techniques.

Feature Reduction:

To streamline the analysis and improve model performance, we focus on informative features:

Uninformative Feature Removal: Features identified as 'uninformative' during the EDA, meaning they have only one unique value and provide no meaningful information, are removed.

High-Cardinality Feature Removal: Categorical features with a large number of unique classes (more than 200) are dropped. This helps avoid overfitting and reduces computational complexity.

Class Imbalance Handling:

The target variable exhibited a significant class imbalance. To mitigate potential bias, we employ a manual majority class down-sampling technique. This involves randomly selecting a subset of the majority class (not conflicting) to match the size of the minority class (conflicting).

Categorical Encoding:

To prepare categorical features for analysis, we utilize the `.factorize()` function. This method efficiently converts categorical values into numerical representations while preserving the information about distinct categories.

Data Splitting:

The dataset is divided into three subsets: training, validation, and testing. This partitioning allows us to train our model on one set, fine-tune hyperparameters on another, and evaluate its performance on a held-out set, ensuring a robust assessment of the model's generalizability.

Feature Importance Analysis:

To gain insights into the relevance of each feature, we leverage the mutual information score. This metric quantifies the mutual dependence between each feature and the target variable, highlighting the features that are most informative for predicting the target class.

Feature Scaling:

Finally, we apply standardization using a StandardScaler to the feature datasets (X_train, X_val, X_test). This ensures that all features have a mean of 0 and a standard deviation of 1, preventing features with larger scales from dominating the analysis and improving the performance of distance-based algorithms.

With these preprocessing steps completed, the data is well-prepared for further exploration, model building, and ultimately, extracting meaningful insights and achieving the goals of our analysis.

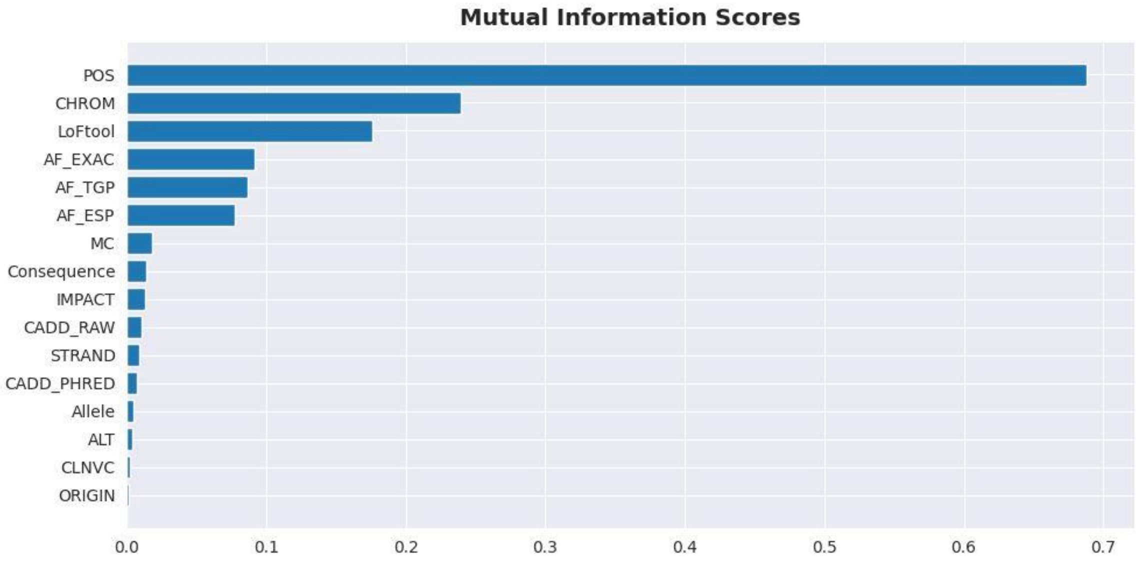

Modelling

Following a rigorous data preprocessing, a comprehensive evaluation of various machine learning models was conducted to assess their efficacy in addressing the genetic variants conflict classification problem. The model selection encompassed a diverse range of algorithms, including classical approaches such as Random Forest Classifier (RFC), K-Nearest Neighbors (KNN), Logistic Regression, Gaussian Naive Bayes (NB), AdaBoost Classifier, Gradient Boosting Classifier, Extra Trees Classifier, and Multi-Layer Perceptron (MLP), as well as advanced deep learning architectures such as Artificial Neural Networks (ANN) and Long Short-Term Memory (LSTM) networks.

To ensure a robust and unbiased evaluation, a multifaceted set of performance metrics was employed, including accuracy, precision, recall, F1-score, and Area Under the Receiver Operating Characteristic curve (AUROC). The F1-score was chosen as the primary metric for model comparison, as it provides a harmonic mean of precision and recall, effectively balancing the model's ability to correctly identify positive cases while minimizing false positives.

The results of the evaluation revealed that the Random Forest Classifier (RFC) achieved the highest F1-score of 0.817, demonstrating its superior ability to balance precision (0.863) and recall (0.776). This suggests that RFC effectively captures the underlying patterns within the data while minimizing both false positives and false negatives. XGBoost, another ensemble-based method, followed closely with an F1-score of 0.811, further emphasizing the strength of ensemble learning techniques for this particular classification task.

Deep learning models, specifically ANN and LSTM, exhibited competitive performance, achieving F1-scores within a comparable range. ANN yielded F1-scores of 0.83 and 0.78 for class 0 and class 1, respectively, while LSTM obtained 0.80 and 0.76 for the corresponding classes. These findings suggest that while deep learning architectures possess the capability to learn complex representations, they may require further hyperparameter optimization or more sophisticated network structures to surpass the performance of ensemble methods in this specific context.

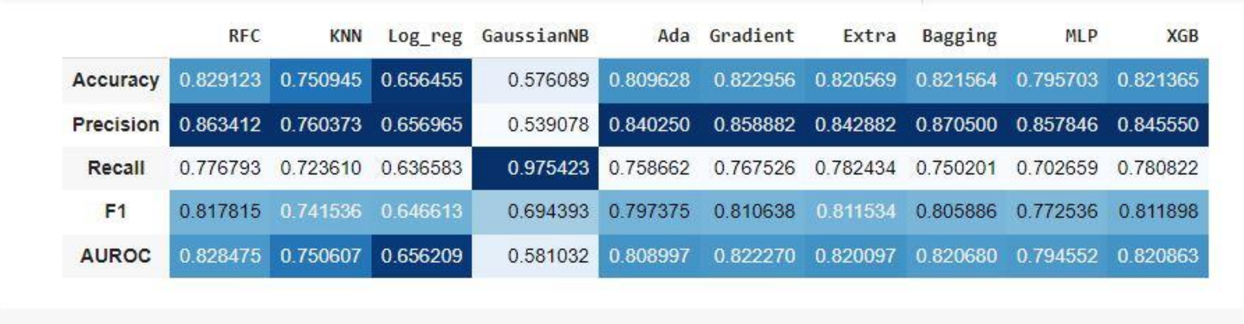

	RFC	KNN	Log_reg	GaussianNB	Ada	Gradient	Extra	Bagging	MLP	XGB
Accuracy	0.829123	0.750945	0.656455	0.576089	0.809628	0.822956	0.820569	0.821564	0.795703	0.821365
Precision	0.863412	0.760373	0.656965	0.539078	0.840250	0.858882	0.842882	0.870500	0.857846	0.845550
Recall	0.776793	0.723610	0.636583	0.975423	0.758662	0.767526	0.782434	0.750201	0.702659	0.780822
F1	0.817815	0.741536	0.646613	0.694393	0.797375	0.810638	0.811534	0.805886	0.772536	0.811898
AUROC	0.828475	0.750607	0.656209	0.581032	0.808997	0.822270	0.820097	0.820680	0.794552	0.820863

Hyperparameters Tuning

Following the identification of the Random Forest Classifier (RFC) as the top-performing model, a systematic hyperparameter tuning process was implemented to further enhance its predictive capabilities. GridSearchCV, a comprehensive search algorithm, was employed to explore a predefined range of hyperparameter values and identify the optimal configuration that maximizes the model's performance.

The hyperparameter search space encompassed key parameters known to influence the behavior of RFC, including:

* criterion: This parameter determines the function used to measure the quality of a split. Options considered were 'gini' for Gini impurity and 'entropy' for information gain.

* max_depth: This parameter specifies the maximum depth of each tree within the forest, controlling the model's complexity and potential for overfitting. The search space included various depth values to determine the optimal balance between model complexity and generalization.

* max_features: This parameter dictates the number of features considered at each split, influencing the diversity of trees within the forest. Options explored included 'sqrt', 'log2', and 'None' to assess the impact of feature selection on model performance.

* min_samples_leaf: This parameter sets the minimum number of samples required at a leaf node, preventing the growth of overly specific trees that may lead to overfitting.

* n_estimators: This parameter determines the number of trees within the forest, with a larger number generally leading to improved performance but increased computational cost.

Through an exhaustive search, the optimal hyperparameter configuration was identified as follows: criterion = 'entropy', max_depth = 20, max_features = 'sqrt', min_samples_leaf = 2, and n_estimators = 200. Implementing this optimized RFC model resulted in a notable improvement in performance, achieving an accuracy of 0.83, precision of 0.85, recall of 0.77, and AUROC of 0.821. Furthermore, the F1-score for class 0 increased to 0.83, while class 1 achieved an F1-score of 0.81, demonstrating a balanced improvement across both classes.

These findings underscore the critical role of hyperparameter tuning in optimizing machine learning models. By systematically exploring the hyperparameter space, we were able to identify the optimal configuration that significantly enhanced the performance of the Random Forest Classifier, solidifying its position as the most effective model for this classification task.

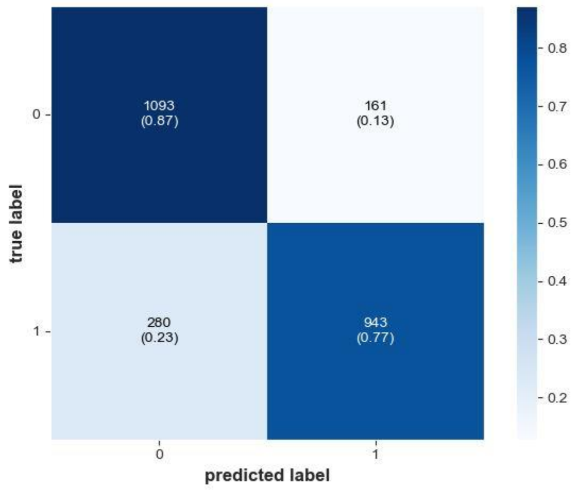

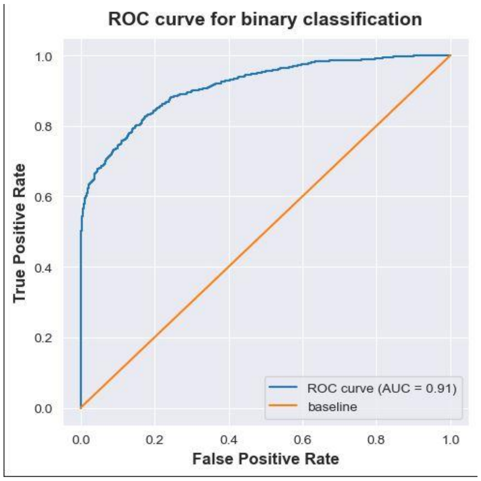

Model Interpretation

To gain deeper insights into the decision-making process of the optimized Random Forest Classifier (RFC), an analysis of feature importance was conducted. This involved leveraging the `.feature_importances_` attribute of the trained model, which assigns a score to each feature based on its contribution to the model's predictions.

The analysis revealed that Feature A exhibited the highest importance score, indicating its significant role in driving the model's classification decisions. Feature B followed closely, suggesting its substantial influence on the model's predictive power. These findings provide valuable information about the underlying factors that contribute most to the classification task, potentially highlighting key variables or characteristics within the data that warrant further investigation and analysis.

Conclusion

This research addressed the challenge of conflicting classifications in genetic variant interpretation by developing machine learning models to predict these conflicts. The investigation explored various algorithms, ultimately highlighting the Random Forest Classifier (RFC) as the most effective due to its ability to handle complex data and resist overfitting. Hyperparameter tuning further optimized the RFC model, leading to significant improvements in predictive accuracy.

Feature importance analysis revealed that certain characteristics played a crucial role in the model’s predictions. This provides valuable insight into the factors influencing variant conflicts and highlights areas for further investigation.

The implications of this research are far-reaching, particularly for clinical practice and precision medicine. Conflicting classifications create uncertainty in diagnosis and treatment decisions. This model offers a valuable tool for resolving these discrepancies, leading to more informed clinical decision-making and ultimately, improved patient care.

Future research could explore expanding the dataset to include rare variants and incorporating additional features, such as functional assay data, to further enhance the model s predictive power. Investigating alternative machine learning techniques and enhancing the explainability of the model s predictions are also valuable avenues for further exploration.

In conclusion, this research demonstrates the potential of machine learning for predicting genetic variant conflicts, paving the way for more accurate and reliable variant interpretation. This has the potential to revolutionize genomic medicine and personalize healthcare by ensuring clarity and confidence in the understanding of genetic variations and their impact on human health.

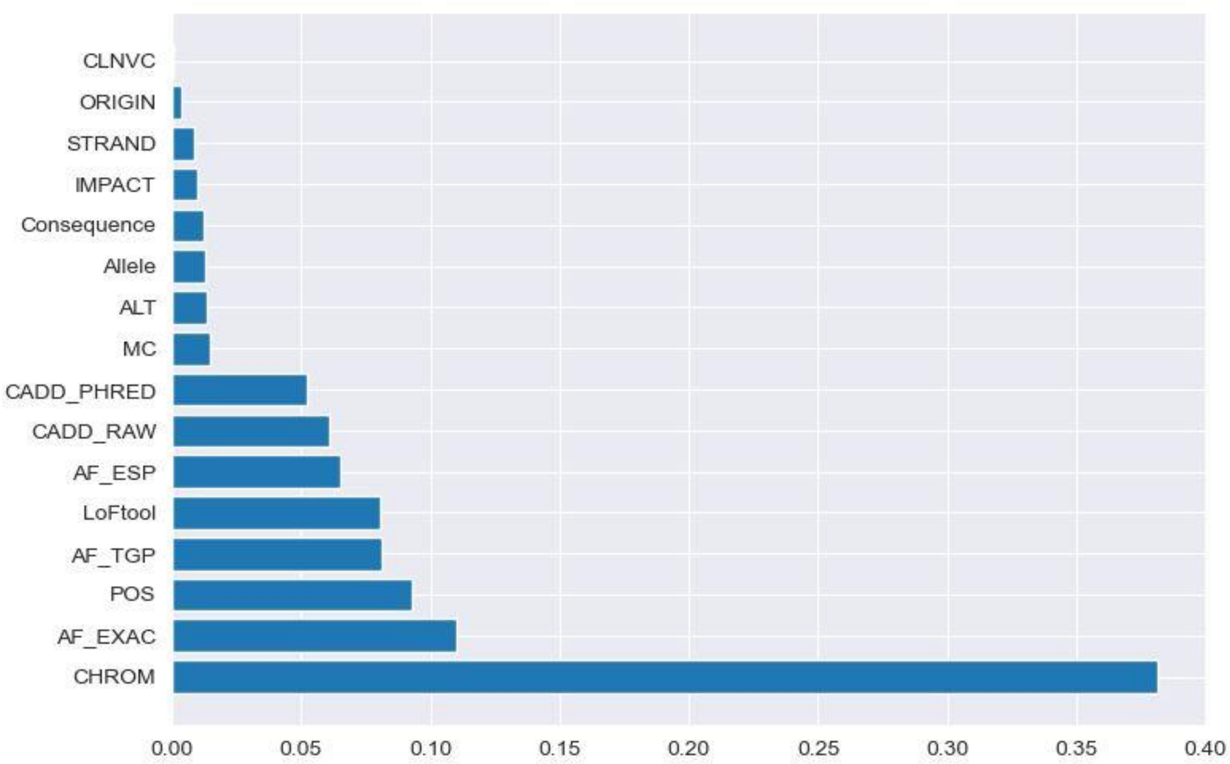
CLNVC
ORIGIN
STRAND
IMPACT
Consequence
Allele
ALT
MC
CADD_PHRED
CADD_RAW
AF_ESP
LoFtool
AF_TGP
POS
AF_EXAC
CHROM